

Akash J

+91-8590875699 | akashj0004@gmail.com | [linkedin.com/akash-aiengineer](https://www.linkedin.com/akash-aiengineer) | github.com/Akash-AldevX

ABOUT ME

Computer Science (AI) graduate with hands-on experience in Azure Data Engineering, ETL pipeline development, and cloud-based data solutions. Skilled in Python, SQL, Azure Data Factory, PostgreSQL, Spark, and Kafka. Experienced in building scalable data pipelines, automating workflows, implementing monitoring solutions and delivering high-quality datasets for analytics and AI/ML applications.

EDUCATION

Dr. MGR Educational and Research Institution

B. Tech in Computer Science and Engineering | CGPA: 7.82/10.0

Chennai, India

Aug. 2021 – May 2025

EXPERIENCE

Associate Data Engineer - Hyderabad

Bexto10x IT Solution — 7 Months

Nov. 2025 – May. 2026

Hyderabad

- Built and maintained scalable ETL pipelines using Azure Data Factory, Python and SQL, improving data processing efficiency.
- Designed data architecture using PostgreSQL, Redis and Azure Blob Storage, supporting reliable batch and Delivered clean, production-ready datasets supporting LLM and ML workflows.
- Automated data Ingestion, validation, Transformation monitoring (Grafana) and deployments using Git & Azure DevOps, ensuring high data quality for AI/ML teams.

Data Science Intern

Digisamaskh IT Company — 4 Months

- Gained experience in data analysis and machine learning using Python and SQL Proficiency in libraries and tools such as Pandas, NumPy, scikit-learn, and data visualization using Power BI.
- Hands-on experience with large-scale data processing on Spark, BigQuery and reproducible analytics workflows.

PROJECTS

CodeLens | Live | CrewAI, LangSmith, FastAPI, Celery, Kafka, Pyspark, Redis, GitHub API, Docker

Mar. 2026

- Engineered end-to-end ETL pipelines using PySpark to ingest, transform, and aggregate GitHub repository metadata, 6-agent CrewAI code review system, agent traces, and review metrics, enabling scalable analytics across large codebases.
- Designed an event-driven data architecture using Kafka, Celery, and Redis for distributed task orchestration, real-time data processing, asynchronous job execution, and high-throughput review workflow management.
- Developed analytics and observability pipelines leveraging LangSmith traces, repository insights, review outcomes, performance monitoring reports, and engineering productivity insights within a Dockerized microservices architecture.

RAG Based PC prediction | Live | LangGraph, LangSmith, RAG, FastAPI, Celery, Redis, Pinecone

Apr. 2026

- Developed a RAG-based recommendation system using FAISS and FastAPI to suggest gaming PC configurations from SKU-level hardware catalogs based on user requirements, budget, and performance preferences.
- Engineered an asynchronous backend using FastAPI, Celery, Redis, async SQLAlchemy, and PostgreSQL for scalable recommendation processing, background indexing of hardware data, user session management and persistent recommendation history. Implemented task queues to handle long-running retrieval and recommendation workflows without blocking API requests.
- Designed and implemented ETL workflows using Python to ingest, clean, transform, and index hardware specification data into PostgreSQL and Pinecone, enabling efficient retrieval and scalable recommendation processing.

Vehicle Insurance Prediction | End-to-End ML Deployment, MongoDB, AWS, Docker, CI/CD

Sep. 2025

- Built an end-to-end data pipeline using Python, Pandas, MongoDB and SQL to ingest, validate, clean, and transform 50K+ vehicle insurance records for machine learning workflows.
- Developed and evaluated multiple classification models for insurance response prediction, implementing feature engineering, experiment tracking and model evaluation to improve prediction performance
- Containerized the application using Docker and automated deployment to Azure App Service through GitHub Actions CI/CD pipelines, reducing manual deployment effort and ensuring reproducible releases.

TECHNICAL SKILLS

Python | SQL | Bash | Azure Data Factory | Azure Data Lake | Azure Blob Storage | Azure DevOps | Apache Spark | Apache Kafka | ETL/ELT Pipelines | PostgreSQL | MySQL | Redis | MongoDB | Docker | Kubernetes | Grafana | Data Warehousing | Data Modeling | CI/CD | Data Governance | GitHub Actions | MLFlow | Model Evaluation Pipelines | Token Optimization | Pinecone | FAISS

CERTIFICATIONS

Udemy | Full Stack Generative and Agentic AI with Python | Hitesh Choudhary & Piyush Garg

Apr. 2026

CampusX | Crew AI | Nitish Singh

Feb. 2026

IBM | Machine Learning with Python

Dec. 2025